

CSc 360
Operating Systems
CPU Scheduling


Wenjun Yang
Summer 2026

Review: process and thread

- Uniprogramming
- Multiprogramming
- Multitasking (multiprocessing)
- Multithreading
- How to handle ***many*** processes/threads?
 - state: ready, running, blocked
 - scheduling (PCB, TCB)


CPU scheduling

- Goal
 - maximize resource utilization
 - CPU, memory, storage
 - improve system responsiveness
- Example
 - CPU burst
 - I/O burst



CPU burst distribution

- Observation
 - many short bursts
 - a few long bursts
- Why is this important?



CPU scheduler

- CPU scheduler
 - short-term scheduling: *who's next?*
- CPU scheduling
 - switch from running to waiting (blocked)
 - switch between running and ready
 - switch from waiting to ready
 - terminate (i.e., leave the system)
- Preemptive vs non-preemptive

CPU dispatcher

- Dispatcher
 - give control to the one selected by scheduler
- Procedures
 - context switching (save old, load new, etc)
 - mode switching (e.g., switch to user mode)
 - start to execute from the newly loaded PC
- Performance measure
 - dispatch latency/overhead

Scheduling criteria

- “Who’s next?”
 - CPU utilization: keep CPU as busy as possible
 - throughput: # of processes done per unit time
 - turnaround time: from start to finish
 - waiting time: time spent in ready state
 - response time: interactive, request-reply
- Goal
 - max {...}, min {...}

Scheduling algorithms

- First come, first serve (FCFS)
 - served by the order of arrival
- Example
 - P1(24), P2(3), P3(3)
 - schedule A (unlucky arrival order)
 - P1, P2, P3
 - schedule B (lucky arrival order)
 - P2, P3, P1

This lecture so far

- CPU scheduling
 - who's who
 - scheduler, dispatcher
 - who's next
 - FCFS (and many more)
- Explore further
 - what's the best way (data structure + algorithm) to implement an FCFS scheduler?




Shortest job first

- SJF: based on the length of *next* CPU burst
 - non-preemptive
 - the job with the smallest burst length scheduled first
 - or preemptive
 - i.e., always the shortest remaining time first
- SJF is optimal in average waiting time
 - reduce the total waiting time for all jobs
 - why *is* SJF optimal?

SJF: example


- Example

- P1: 0 (arrival time); 7 (burst time)
- P2: 2; 4
- P3: 4; 1
- P4: 5; 4



- Non-preemptive⁰

- P1, P3, P2, P4



- Preemptive

- P1, P2, P3, P2, P4, P1



CSc 360

11

Shortest Remaining Time First (SRTF)

Q: turnaround time?

SJF: more

- Determine the *next* burst length
 - how to predict the future?
 - if history is of any indication ...
 - use the last burst length
 - use the average so far
 - use the moving average
 - use the weighted moving average
 - the exponentially weighted moving average
 - i.e., $\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$

CSc 360

12

Q: storage requirement?


EWMA: example

- Exponentially weighted moving average

- $\tau_0 = 10$

- $\alpha = 0.5$

- normally (0,1)



CPU burst (t_i)	6	4	6	4	13	13	13	...	
"guess" (τ_i)	10	8	6	6	5	9	11	12	...

CSc 360

Q: when $\alpha=0$ or 1?

Priority scheduling

- Priority
 - the job with the highest priority scheduled first
 - SJF: shorter CPU burst, higher priority
 - FCFS: arrival earlier, higher priority
 - static priority: starvation
 - e.g., SJF
 - dynamic priority
 - e.g., aging
- Non-preemptive vs preemptive

Round-robin scheduling


- Discrete processor sharing
 - CPU time quantum
 - usually 10~100 ms
 - for a process
 - either yield after a CPU burst
 - or be preempted after using up a time quantum
 - a FIFO queue
 - all ready processes
- Weighted round-robin

RR: example

- Example
 - P1: 0 (arrival time); 7 (burst time)
 - P2: 2; 4
 - P3: 4; 1
 - P4: 5; 4
- Time quantum
 - e.g., 1 quantum = 1 time unit
 - how about 1 quantum = 4 time units

Time quantum

- Large quantum
 - => FCFS
- Small quantum
 - better responsiveness
 - be aware of overhead
 - context switching
 - “80%” rule



process	time
P_1	6
P_2	3
P_3	1
P_4	7

The 2nd half of this lecture


- Scheduling algorithms
 - FCFS
 - SJF, priority, RR
 - preemptive and non-preemptive
- Explore further
 - evaluate average waiting time, average turnaround time per unit job for these algorithms
 - Q's on slides and in the textbook

Next in the lecture

- Next in the lecture
 - more on scheduling


Multi-queue scheduling

- “No one fits all”
- Multi-queue approach
 - foreground queue
 - e.g., RR; better fairness
 - background queue
 - e.g., FCFS; more efficient
- Inter-queue scheduling
 - priority, time sharing (e.g., “80% rule”)



Multi-queue with feedback

- Multi-queue
 - number of queues
 - queuing algorithm for each queue
- Multi-queue with feedback
 - promote jobs
 - demote jobs




Multi-processor queuing

- Load balance between processors
 - cooperation and communication
- Asymmetric multiprocessing
 - one master scheduler
- Symmetric multiprocessing
 - cooperative schedulers
 - processor affinity: try to stick with one
 - load balancing: push or pull migration

More on scheduling

- Thread scheduling
 - local: user $\rightarrow$ kernel thread
 - e.g., within a process
 - global: kernel thread $\rightarrow$ CPU
 - e.g., across the system
- Algorithm evaluation
 - queuing analysis
 - Little's law: $n = \lambda * W$
 - simulation



Pthread scheduling

- `pthread_attr_setschedpolicy ()`;
 - regular, non-realtime (nice-able)
 - realtime, round-robin (preemptive, privileged)
 - realtime, FCFS (non-preemptive, privileged)
- `pthread_attr_setschedparam ()`;
- `pthread_attr_setscope ()`;
 - scheduling within a process
 - scheduling for the entire system

This lecture finally

- More on scheduling
 - multi-queue scheduling
 - multiprocessor scheduling
 - scheduling evaluation
 - Little's law, simulation
- Explore further
 - list process priority: `/usr/bin/top`
 - change priority: `/bin/nice`

Next lecture

- Process synchronization
 - read OSC7 Chapter 6 (or OSC6 Chapter 7)